

```
In [1]: ▶ import pandas as pd
from sklearn.preprocessing import OneHotEncoder
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
```

```
In [2]: ▶ my_data = pd.read_csv('bmi.csv')
my_data
```

```
Out[2]:
```

	Age	Height	Weight	Bmi	BmiClass
0	61	1.85	109.30	31.935720	Obese Class 1
1	60	1.71	79.02	27.023700	Overweight
2	60	1.55	74.70	31.092612	Obese Class 1
3	60	1.46	35.90	16.841809	Underweight
4	60	1.58	97.10	38.896010	Obese Class 2
...	...	...	...	...	...
736	34	1.86	95.70	27.662157	Overweight
737	44	1.91	106.90	29.302925	Overweight
738	25	1.82	88.40	26.687598	Overweight
739	35	1.88	98.50	27.868945	Overweight
740	45	1.93	109.90	29.504148	Overweight

741 rows × 5 columns

```
In [3]: ▶ X = my_data.drop(columns=['Bmi'])
y = my_data['Bmi']
print(my_data.dtypes)
```

```
Age          int64
Height       float64
Weight       float64
Bmi          float64
BmiClass     object
dtype: object
```

```
In [4]: data_encoded = pd.get_dummies(my_data, columns=['BmiClass'])
print(data_encoded.head())
onehot_encoder = OneHotEncoder()
encoded_column = onehot_encoder.fit_transform(my_data['BmiClass'].values.reshape(-1, 1))
encoded_df = pd.DataFrame(encoded_column.toarray(), columns=onehot_encoder.get_feature_names_out())
data_encoded = pd.concat([my_data.drop(columns=['BmiClass']), encoded_df], axis=1)
print(data_encoded.head())
```

	Age	Height	Weight	Bmi	BmiClass_Normal	Weight \
0	61	1.85	109.30	31.935720		False
1	60	1.71	79.02	27.023700		False
2	60	1.55	74.70	31.092612		False
3	60	1.46	35.90	16.841809		False
4	60	1.58	97.10	38.896010		False

	BmiClass_Obese Class 1	BmiClass_Obese Class 2	BmiClass_Obese Class 3
0	True	False	False
1	False	False	False
2	True	False	False
3	False	False	False
4	False	True	False

	BmiClass_Overweight	BmiClass_Underweight
0	False	False
1	True	False
2	False	False
3	False	True
4	False	False

	Age	Height	Weight	Bmi	BmiClass_Normal	Weight \
0	61	1.85	109.30	31.935720		0.0
1	60	1.71	79.02	27.023700		0.0
2	60	1.55	74.70	31.092612		0.0
3	60	1.46	35.90	16.841809		0.0
4	60	1.58	97.10	38.896010		0.0

	BmiClass_Obese Class 1	BmiClass_Obese Class 2	BmiClass_Obese Class 3
0	1.0	0.0	0.0
1	0.0	0.0	0.0
2	1.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	1.0	0.0

	BmiClass_Overweight	BmiClass_Underweight
0	0.0	0.0
1	1.0	0.0
2	0.0	0.0
3	0.0	1.0
4	0.0	0.0

```
In [5]: X = data_encoded.drop(columns=['Bmi'])
y = data_encoded['Bmi']
X
```

Out[5]:

	Age	Height	Weight	BmiClass_Normal Weight	BmiClass_Obese Class 1	BmiClass_Obese Class 2	BmiClass_
0	61	1.85	109.30	0.0	1.0	0.0	
1	60	1.71	79.02	0.0	0.0	0.0	
2	60	1.55	74.70	0.0	1.0	0.0	
3	60	1.46	35.90	0.0	0.0	0.0	
4	60	1.58	97.10	0.0	0.0	1.0	
...	...	...	...	...	...	...	
736	34	1.86	95.70	0.0	0.0	0.0	
737	44	1.91	106.90	0.0	0.0	0.0	
738	25	1.82	88.40	0.0	0.0	0.0	
739	35	1.88	98.50	0.0	0.0	0.0	
740	45	1.93	109.90	0.0	0.0	0.0	

741 rows × 9 columns



```
In [6]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r

model = LinearRegression()

model.fit(X_train, y_train)

y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error:", mse)
```

Mean Squared Error: 0.7690052384587545

```
In [7]: r_squared = model.score(X_test, y_test)
print("R-squared:", r_squared)
```

R-squared: 0.9892416129298092

## FINE TUNING

```
In [8]: param_grid = {  
        'copy_X': [True, False],  
        'fit_intercept': [True, False],  
        'n_jobs': [True, False],  
        'positive': [True, False]  
    }
```

```
In [9]: grid_search = GridSearchCV(model, param_grid, cv=5)  
grid_search.fit(X_train, y_train)
```

```
Out[9]: GridSearchCV(cv=5, estimator=LinearRegression(),  
                    param_grid={'copy_X': [True, False],  
                                'fit_intercept': [True, False],  
                                'n_jobs': [True, False], 'positive': [True, Fals  
e]})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with [nbviewer.org](https://nbviewer.org).**

```
In [10]: best_params = grid_search.best_params_  
best_model = grid_search.best_estimator_
```

```
In [11]: best_model.fit(X_train, y_train)  
y_pred_best = best_model.predict(X_test)  
mse_best = mean_squared_error(y_test, y_pred_best)  
print("Fine-Tuned Mean Squared Error:", mse_best)
```

Fine-Tuned Mean Squared Error: 0.7690052384586076

```
In [12]: ▶ plt.scatter(y_test, y_pred_best)
plt.xlabel("Actual Values")
plt.ylabel("Predicted Values")
plt.title("Predictions vs Actuals")
plt.show()
```

